

MD ZAID

Agentic AI Engineer · Agent Harness & Orchestration · Tool Calling, Planning & Memory · RAG / Vector Retrieval · Agent Evals

+91 93219 87363 | mdzaid19433@gmail.com | linkedin.com/in/mdzaid433 | Mumbai / Gurgaon, India

PROFESSIONAL SUMMARY

Agentic AI Engineer with 3+ years of experience at Jio Platforms (Reliance Industries), specializing in production-grade AI agents, multi-agent systems, and **RAG** architectures. Strong expertise in agent orchestration, LLM evaluation, **guardrails**, memory, observability, and production reliability, with hands-on experience across **LangGraph**, **Claude**, **Azure Databricks**, **MLflow**, **DSPy**, and **Databricks Vector Search**. Experienced in owning AI systems end-to-end, designing robust agent workflows, implementing evaluation and feedback loops, and deploying scalable, reliable enterprise GenAI solutions.

TECHNICAL SKILLS

Agentic AI	Agent design & orchestration, tool calling, MCP , ReAct planning, LangGraph StateGraph, multi-agent (worker + auditor), short & long-term memory, human-in-the-loop, guardrails & policy enforcement
Agent Harness & Eval	MLflow (tracing, prompt registry, versioning), Langfuse, LLM-as-Judge , SME-aligned eval datasets, automated prompt optimization (DSPy), groundedness / faithfulness / citation metrics, production monitoring
RAG & Retrieval	Hybrid dense + sparse retrieval, runtime query routing (SQL ↔ vector), chunking strategies, embedding lifecycle & re-indexing, Databricks Vector Search index tuning, context assembly & prompt orchestration
LLM & Frameworks	LangChain, LangGraph, DSPy, LiteLLM, Claude, OpenAI, HuggingFace, FastAPI, LoRA / QLoRA
Platform & Languages	Azure Databricks , Unity Catalog, Lakebase, Docker, Kubernetes, Redis, DAB · Python, Java, JavaScript, SQL

EXPERIENCE

AI / Agentic AI Engineer

Aug 2023 – Present

Jio Platforms Ltd. (Reliance Industries), Mumbai / Gurgaon, India

Email Copilot & Auditor Agent | Jio KM Portal (Multi-agent LangGraph · Claude Sonnet · Databricks · 11 business portals, 7.8M requests/day)

- Problem statement:** 10,000+ customer emails a day across 11 business portals were triaged and drafted manually, with no automated check on accuracy, tone or PII exposure.
- Architected the agent harness as a multi-node **LangGraph** conditional state graph: typed graph state, conditional planning and routing, **MCP** tool calls into enterprise systems, retry/fallback branches, and a 0-10 confidence score gating autonomous send vs. human escalation.
- Built the drafting pipeline as graph stages: queue classification, intent detection, ABSA sentiment and **RAG**-grounded drafting via **DSPy** ChainOfThought, with PII masking (Aadhaar, PAN, mobile, DOB) enforced as a hard **policy guardrail** before any draft exits the graph.
- Added an independent **Auditor Agent** that scores every draft against **9 guardrails** (≥80% to pass), runs a max-3-attempt HITL revision loop and permanently rejects non-compliant drafts.
- Cut latency and token cost with a T2R (Travel-to-Resolution) fast path that drafts from pre-fetched **MCP** tool responses, bypassing full retrieval.
- Impact:** 10,000+ emails/day resolved autonomously across all 11 portals, 95% reduction in manual handling effort, zero non-compliant drafts released.
- Tools & Tech:** LangGraph (StateGraph, conditional routing), MCP, DSPy, Claude Sonnet, RAG, Databricks, Python, FastAPI

REIMS | Enterprise Plant SOP Agent (Azure Databricks · Claude Sonnet · Databricks Vector Search · multi-site deployment)

- Problem statement:** plant teams needed citable answers from thousands of pages of structured and unstructured SOPs across multiple sites, where any cross-site data exposure is a compliance failure.
- Drove the build from solution design through UAT sign-off and production deployment, covering tool design, retrieval layer, memory and evaluation.
- Implemented dual memory: short-term per-thread state plus long-term user memory persisted in **Lakebase**, exposed as read / write / delete tools the agent invokes autonomously.
- Enforced site isolation inside the retrieval layer, filtering every query by the authenticated user's authorized sites rather than by prompt instruction.
- Owned retrieval quality: chunking strategy, embedding lifecycle, re-indexing and **Databricks Vector Search** index tuning across structured and unstructured SOP content.
- Built a self-improving eval harness: thumbs up/down → **MLflow** traces → 500-example SME-reviewed eval set → automated prompt optimization (**DSPy**) scored by SME-aligned LLM judges → winning prompt promoted via MLflow versioning, no redeploy.
- Impact:** p95 latency under 4s, groundedness ≥0.85, citation completeness ≥0.95, zero cross-site data leaks in production.
- Tools & Tech:** Azure Databricks, Databricks Vector Search, Lakebase, MLflow, DSPy, LLM-as-Judge, Claude Sonnet, Python

Vendor Agreement Agent | Reliance Life Sciences (Internal legal teams · Hybrid retrieval over structured DB + contract documents)

- Problem statement:** vendor-agreement queries required cross-referencing structured contract metadata in databases against clause text inside the agreements; neither SQL nor dense retrieval alone covered both.
- Built a hybrid-retrieval agent that classifies each query at runtime and routes it between **SQL** over structured contract data and **dense document retrieval**.
- Led extraction and modelling of contract key-value data, letting the agent fuse structured lookups with unstructured clause retrieval into one grounded answer.
- Impact:** 40%+ improvement in answer accuracy with reduced query latency, replacing the manual cross-referencing workflow for legal teams.
- Tools & Tech:** LangChain / LangGraph, SQL (structured contract DB), dense vector retrieval, Claude, Python

EDUCATION

B.Tech, Computer Science & Biosciences

2019 – 2023

IIITD (Indraprastha Institute of Information Technology Delhi), New Delhi, India